

Internship_CodeB_week 4

April 19, 2025

1 Phishing Website Detection

- Name : Gaurav Vijay Jadhav
- github : [https://github.com/jadhavgaurav/CodeB_Internship_Project]

2 Week 4 Submission

```
[25]: # Import Necessary Libraries

import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import warnings
warnings.filterwarnings('ignore')
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler, LabelEncoder
```

```
[26]: # import dataset

data_url = 'https://raw.githubusercontent.com/jadhavgaurav/
↳CodeB_Internship_Project/refs/heads/main/dataset_phishing.csv'

df = pd.read_csv(data_url)

df.sample(frac = 1)
```

```
[26]:
```

	url	length_url	\
9990	https://beta.znipe.tv/	22	
10367	https://www.armeriaegara.com/	29	
7134	http://www.techsupportalert.com/content/best-f...	74	
7535	http://www.forlocations.com/savealot	36	
3676	http://www.howdesign.com/editors-picks/10-eye-...	80	
...	
3545	http://houseoftiresbcs.com/Adobe/css/XML/PDF/	45	
9653	http://www.yilport.n.nu/	24	
1533	http://smsenligne.myfreesites.net/	34	

11223	http://www.waymarking.com	25
5677	https://www.osapublishing.org/abstract.cfm?URI...	61

	length_hostname	ip	nb_dots	nb_hyphens	nb_at	nb_qm	nb_and	nb_or	\
9990	13	0	2	0	0	0	0	0	
10367	20	0	2	0	0	0	0	0	
7134	24	0	3	4	0	0	0	0	
7535	20	0	2	0	0	0	0	0	
3676	17	0	2	6	0	0	0	0	
...	
3545	19	0	1	0	0	0	0	0	
9653	16	0	3	0	0	0	0	0	
1533	26	0	2	0	0	0	0	0	
11223	18	0	2	0	0	0	0	0	
5677	21	0	3	3	0	1	0	0	

	domain_in_title	domain_with_copyright	whois_registered_domain	\
9990	0	0	0	
10367	1	1	0	
7134	1	1	0	
7535	1	0	0	
3676	1	0	0	
...	
3545	1	0	0	
9653	0	0	0	
1533	1	0	0	
11223	0	1	0	
5677	1	1	0	

	domain_registration_length	domain_age	web_traffic	dns_record	\
9990	173	1654	4127669	0	
10367	171	4578	0	0	
7134	79	7226	52982	0	
7535	14	3639	246250	0	
3676	218	8915	566894	0	
...	
3545	53	1407	0	0	
9653	2692	-1	0	0	
1533	242	1950	341948	0	
11223	67	5776	197005	0	
5677	1697	2320	35081	0	

	google_index	page_rank	status
9990	0	3	legitimate
10367	0	2	legitimate
7134	0	5	legitimate
7535	0	3	legitimate

3676	1	5	legitimate
...
3545	1	0	phishing
9653	0	3	legitimate
1533	1	1	phishing
11223	0	5	legitimate
5677	0	5	legitimate

[11430 rows x 89 columns]

3 Data Cleaning Report Phishing Website Detection

3.1 Dataset Overview

- **Total Records:** 11,430
- **Total Features (excluding target):** 87
- **Target Variable:** status
 - 0: Legitimate
 - 1: Phishing
- **Data Types:**
 - Numerical (int64/float64): 87
 - Categorical/Object: 1 (url)

3.1.1 Target Column

3.1.2 status

- **Description:** Binary label indicating if the website is phishing (1) or legitimate (0).
 - **Relevance:** This is the variable to be predicted by the classification model.
-

```
[29]: # Replace 'Legitimate' with 0 and 'Phishing' with 1 in the 'status' column
df['status'] = df['status'].map({'legitimate':0, 'phishing':1})

print(df['status'].value_counts())
```

```
status
0    5715
1    5715
Name: count, dtype: int64
```


```
[30]: # Basic Info About Target Column and Visualize Target Distribution (Bar Plot)

# Check class distribution

sns.countplot(data=df, x='status', palette='Set2')
plt.title("Target Variable Distribution (Phishing vs Legitimate)")
```

```
plt.show()

print(df['status'].value_counts())
```



```
status
0    5715
1    5715
Name: count, dtype: int64
```

```
[31]: numeric_features = df.select_dtypes(include=['int64', 'float64']).columns.
      ↪ tolist()
      categorical_features = df.select_dtypes(include='object').columns.tolist()

      print("Numeric Features:", numeric_features)
      print("Categorical Features:", categorical_features)
```

```
Numeric Features: ['length_url', 'length_hostname', 'ip', 'nb_dots',
'nb_hyphens', 'nb_at', 'nb_qm', 'nb_and', 'nb_or', 'nb_eq', 'nb_underscore',
'nb_tilde', 'nb_percent', 'nb_slash', 'nb_star', 'nb_colon', 'nb_comma',
'nb_semicolumn', 'nb_dollar', 'nb_space', 'nb_www', 'nb_com', 'nb_dslash',
'http_in_path', 'https_token', 'ratio_digits_url', 'ratio_digits_host',
```

```
'punycode', 'port', 'tld_in_path', 'tld_in_subdomain', 'abnormal_subdomain',
'nb_subdomains', 'prefix_suffix', 'random_domain', 'shortening_service',
'path_extension', 'nb_redirection', 'nb_external_redirection',
'length_words_raw', 'char_repeat', 'shortest_words_raw', 'shortest_word_host',
'shortest_word_path', 'longest_words_raw', 'longest_word_host',
'longest_word_path', 'avg_words_raw', 'avg_word_host', 'avg_word_path',
'phish_hints', 'domain_in_brand', 'brand_in_subdomain', 'brand_in_path',
'suspicious_tld', 'statistical_report', 'nb_hyperlinks', 'ratio_intHyperlinks',
'ratio_extHyperlinks', 'ratio_nullHyperlinks', 'nb_extCSS',
'ratio_intRedirection', 'ratio_extRedirection', 'ratio_intErrors',
'ratio_extErrors', 'login_form', 'external_favicon', 'links_in_tags',
'submit_email', 'ratio_intMedia', 'ratio_extMedia', 'sfh', 'iframe',
'popup_window', 'safe_anchor', 'onmouseover', 'right_click', 'empty_title',
'domain_in_title', 'domain_with_copyright', 'whois_registered_domain',
'domain_registration_length', 'domain_age', 'web_traffic', 'dns_record',
'google_index', 'page_rank', 'status']
Categorical Features: ['url']
```

```
[32]: # Dropping the 'url' column
      # The 'url' column is not useful for training the machine learning model.


      df.drop(columns=['url'], inplace=True)
```

4 Feature Selection Report

Step 1: Correlation Analysis

Remove features that are highly correlated with each other (e.g., correlation > 0.9 or < -0.9) to reduce multicollinearity.

```
[33]: # Step 1: Compute correlation matrix
      corr_matrix = df.drop('status', axis=1).corr().abs() # Exclude target column
      plt.figure(figsize=(22, 18))
      sns.heatmap(corr_matrix, annot=False, cmap='coolwarm', linewidths=0.5)
      plt.title("Correlation Matrix")
      plt.show()
```



- The correlation heatmap was generated to visually inspect multicollinearity between features.
- Correlation threshold used: 0.90

Heatmap legend:

Red diagonal = perfect correlation (with itself)

Light blue = weak or no correlation

Orange/red= strong correlation

```
[34]: # Step 2: Select upper triangle of correlation matrix
upper = corr_matrix.where(np.triu(np.ones(corr_matrix.shape), k=1).astype(bool))

# Step 3: Find features with correlation > 0.9
to_drop = [column for column in upper.columns if any(upper[column] > 0.9)]
print(f"Highly correlated features to drop (corr > 0.9):\n{to_drop}")
```

```
# Step 4: Drop the features from the dataset
df_reduced = df.drop(columns=to_drop)
print(f"\nShape before dropping: {df.shape}")
print(f"Shape after dropping: {df_reduced.shape}")
```

Highly correlated features to drop (corr > 0.9):
['nb_eq', 'longest_word_path']

Shape before dropping: (11430, 88)

Shape after dropping: (11430, 86)

- Computed the correlation matrix (Pearson correlation).
- Identified pairs of features with absolute correlation > 0.90.
- From each such pair, one feature was dropped to reduce redundancy.

Dropped Features:

- Based on correlation > 0.90, the following features were removed:

'nb_eq'

'longest_word_path'

- These features were highly correlated with other features carrying similar information.

```
[35]: df_reduced.drop(columns=['avg_word_host'], inplace=True) # Drop avg_word_host_
      ↪ column as per VIF analysis
```

4.0.1 2: Feature Selection using ANOVA F-test (f_classif)

```
[36]: from sklearn.feature_selection import SelectKBest, f_classif

X = df_reduced.drop(columns=['status'])
y = df_reduced['status']

# Apply ANOVA F-test
selector = SelectKBest(score_func=f_classif, k=30) # Select top 20 features
X_kbest = selector.fit_transform(X, y)

# Get selected feature names
selected_features_f_classif = X.columns[selector.get_support()]
print("Top 30 Features selected using f_classif:")
print(selected_features_f_classif)
```

Top 30 Features selected using f_classif:

```
Index(['length_url', 'length_hostname', 'ip', 'nb_dots', 'nb_qm', 'nb_and',
      'nb_slash', 'nb_www', 'ratio_digits_url', 'ratio_digits_host',
      'tld_in_subdomain', 'prefix_suffix', 'length_words_raw',
      'shortest_word_host', 'longest_words_raw', 'avg_words_raw',
```

```

    'avg_word_path', 'phish_hints', 'nb_hyperlinks', 'ratio_intHyperlinks',
    'links_in_tags', 'ratio_intMedia', 'safe_anchor', 'empty_title',
    'domain_in_title', 'domain_with_copyright',
    'domain_registration_length', 'domain_age', 'google_index',
    'page_rank'],
    dtype='object')

```

4.0.2 3: Random Forest Feature Importance

```

[37]: from sklearn.ensemble import RandomForestClassifier
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns


# Load dataset (assuming df is already preprocessed and target is separated)
X = df_reduced.drop('status', axis=1)
y = df_reduced['status']

# Train Random Forest
rf = RandomForestClassifier(n_estimators=100, random_state=42)
rf.fit(X, y)

# Get feature importances
importances = pd.Series(rf.feature_importances_, index=X.columns)
top_30_features = importances.sort_values(ascending=False).head(30)

# Plot
plt.figure(figsize=(12, 8))
sns.barplot(x=top_30_features.values, y=top_30_features.index,
            palette='viridis')
plt.title('Top 30 Features by Random Forest Importance')
plt.xlabel('Importance Score')
plt.ylabel('Feature')
plt.tight_layout()
plt.show()

```

4.0.3 4: Apply RFE (Recursive Feature Elimination)

```
[38]: from sklearn.linear_model import LogisticRegression
from sklearn.feature_selection import RFE

# Use top 40 features for RFE
X_top30 = X[top_30_features.index]

# Apply RFE with Logistic Regression
lr = LogisticRegression(solver='liblinear', random_state=42)
rfe = RFE(estimator=lr, n_features_to_select=20)
rfe.fit(X_top30, y)

# Get selected feature names
selected_features_rfe = X_top30.columns[rfe.support_]
print("Top 20 features selected by RFE:\n")
print(selected_features_rfe)
```

Top 20 features selected by RFE:

```
Index(['google_index', 'page_rank', 'nb_www', 'ratio_extHyperlinks',
      'phish_hints', 'ratio_intHyperlinks', 'ratio_digits_url',
      'avg_word_path', 'longest_words_raw', 'length_words_raw',
      'ratio_extRedirection', 'domain_in_title', 'nb_dots', 'ip',
```

```

        'ratio_digits_host', 'shortest_word_host', 'avg_words_raw',
        'nb_hyphens', 'nb_qm', 'shortest_word_path'],
        dtype='object')

```

4.0.4 Final Selected Features from

selected_features_rfe → top 20 features from RFE on top 30 RF features

selected_features_f_classif → top 30 features from f_classif

```

[39]: # Convert both to sets
rfe_features_set = set(selected_features_rfe)
f_classif_features_set = set(selected_features_f_classif)

# Take intersection
final_selected_features = list(rfe_features_set.union(f_classif_features_set))

print("Final Selected Features (Intersection of RFE and f_classif):")
print(final_selected_features)
print(f"Number of final selected features: {len(final_selected_features)}")

```

```

Final Selected Features (Intersection of RFE and f_classif):
['prefix_suffix', 'nb_hyphens', 'domain_registration_length',
'ratio_digits_host', 'ratio_intHyperlinks', 'avg_word_path', 'ratio_intMedia',
'ip', 'length_hostname', 'length_words_raw', 'nb_www', 'empty_title',
'ratio_extRedirection', 'nb_dots', 'nb_slash', 'tld_in_subdomain',
'nb_hyperlinks', 'google_index', 'nb_and', 'length_url', 'domain_in_title',
'shortest_word_path', 'longest_words_raw', 'ratio_digits_url',
'shortest_word_host', 'links_in_tags', 'safe_anchor', 'page_rank', 'domain_age',
'phish_hints', 'domain_with_copyright', 'avg_words_raw', 'ratio_extHyperlinks',
'nb_qm']
Number of final selected features: 34

```

```

[40]: from statsmodels.stats.outliers_influence import variance_inflation_factor
import pandas as pd

# Subset the dataframe to final selected features
X_vif = df_reduced[final_selected_features]

# X_vif = X_vif.drop(columns=['avg_word_host'])

# Compute VIF
vif_data = pd.DataFrame()
vif_data["Feature"] = X_vif.columns
vif_data["VIF"] = [variance_inflation_factor(X_vif.values, i) for i in
    ↪range(X_vif.shape[1])]

# Sort VIF descending
vif_data = vif_data.sort_values(by="VIF", ascending=False)

```

```
print("VIF for Final Selected Features:")
print(vif_data)
```

VIF for Final Selected Features:

	Feature	VIF
9	length_words_raw	26.646784
31	avg_words_raw	22.677655
19	length_url	21.356093
4	ratio_intHyperlinks	17.695732
14	nb_slash	16.004551
13	nb_dots	11.074524
8	length_hostname	10.018208
22	longest_words_raw	8.455121
5	avg_word_path	8.351384
25	links_in_tags	7.433357
20	domain_in_title	5.553849
27	page_rank	5.356246
24	shortest_word_host	4.893766
23	ratio_digits_url	4.508979
28	domain_age	4.483027
32	ratio_extHyperlinks	4.213635
7	ip	3.718532
17	google_index	3.634672
18	nb_and	2.981160
6	ratio_intMedia	2.974564
1	nb_hyphens	2.916684
21	shortest_word_path	2.902547
26	safe_anchor	2.861841
10	nb_www	2.772176
11	empty_title	2.450279
33	nb_qm	2.128720
30	domain_with_copyright	2.101064
15	tld_in_subdomain	1.972449
29	phish_hints	1.833878
3	ratio_digits_host	1.726988
0	prefix_suffix	1.718168
2	domain_registration_length	1.603927
16	nb_hyperlinks	1.593017
12	ratio_extRedirection	1.540672

```
[41]: features_to_drop_vif = [
        'length_words_raw',
        'avg_words_raw',
        'length_url',
        'ratio_intHyperlinks',
        'nb_slash',
```

```
    'nb_dots',  
]
```

```
[42]: # Final Set of Features After VIF Cleaning
```

```
final_features_vif = list(set(final_selected_features) -  
    ↪ set(features_to_drop_vif))  
print(f"Number of final features after VIF cleaning: {len(final_features_vif)}")  
print("Final Features After VIF Cleaning:")  
final_features_vif
```

Number of final features after VIF cleaning: 28

Final Features After VIF Cleaning:

```
[42]: ['prefix_suffix',  
      'nb_hyphens',  
      'domain_registration_length',  
      'ratio_digits_host',  
      'avg_word_path',  
      'ratio_intMedia',  
      'ip',  
      'length_hostname',  
      'nb_www',  
      'empty_title',  
      'ratio_extRedirection',  
      'tld_in_subdomain',  
      'nb_hyperlinks',  
      'google_index',  
      'nb_and',  
      'domain_in_title',  
      'shortest_word_path',  
      'longest_words_raw',  
      'ratio_digits_url',  
      'shortest_word_host',  
      'links_in_tags',  
      'safe_anchor',  
      'page_rank',  
      'domain_age',  
      'phish_hints',  
      'ratio_extHyperlinks',  
      'domain_with_copyright',  
      'nb_qm']
```

4.1 Applied Steps for Feature Selection Process:

4.1.1 1. Correlation Analysis

- Removed highly correlated features ($\text{corr} > 0.9$)
 - **Dropped:** 'nb_eq', 'longest_word_path'
 - **Reduced from 88 to 86 features**
-

4.1.2 2. ANOVA (f_classif)

- Selected **top 30 features** based on **univariate F-test**
 - Suitable for **numerical features** with **categorical target**
-

4.1.3 3. Random Forest Feature Importance

- Trained a **Random Forest Classifier**
 - Retrieved **top 30 features** using `feature_importances_`
-

4.1.4 4. Recursive Feature Elimination (RFE)

- Applied **RFE** with Random Forest as estimator
 - Selected another **top 30 important features**
-

4.1.5 5. Feature Union

- Took **intersection** of `f_classif_features_set` & `rfe_features_set`
 - Created a **robust final feature set** using two strong methods
-

4.1.6 6. Variance Inflation Factor (VIF)

- Evaluated multicollinearity in final selected features
- Dropped 6 features with $\text{VIF} > 10$ to avoid redundancy

5 Feature Engineering

```
[43]: # 1. URL Complexity Score
# Combines counts of common "suspicious" tokens into a single indicator.
# Phishing URLs often cram many special characters (www, -, ?, &) to obfuscate_
# their true destination.

df_reduced['url_complexity'] = (
    df_reduced['nb_www']
    + df_reduced['nb_hyphens']
```

```

+ df_reduced['nb_qm']
+ df_reduced['nb_and']
)

```

[44]: *# 2. Tag-to-Link Ratio*
Measures the density of "hidden" tags relative to visible hyperlinks.
Fake pages load script/link tags disproportionately to real hyperlinks-high
↪ratios indicate suspicious embedding.

```

df_reduced['tag_to_link_ratio'] = df_reduced['links_in_tags'] /
↪(df_reduced['nb_hyperlinks'] + 1)

```

[45]: *# 3. Domain Numeric Intensity*
Scales the digit-density in the hostname by domain age (older domains with
↪many digits are rarer).
Young domains with a high digit ratio are more likely auto-generated by
↪attackers; multiplying by domain_age highlights this risk.

```

df_reduced['domain_numeric_intensity'] = df_reduced['ratio_digits_host'] *
↪df_reduced['domain_age']

```

[46]: *# 4. Path Word Complexity*
Captures both the average word length and the longest word in the URL path.
Extremely long or complex path segments often appear in phishing payload URLs-
this combines average and maximum word length in the path.

```

df_reduced['path_word_complexity'] = df_reduced['avg_word_path'] *
↪df_reduced['longest_words_raw']

```

[47]: *# Drop 5 low-importance/redundant features*

```

features_to_drop = [
    'domain_with_copyright',
    'ratio_intMedia',
    'google_index',
    'page_rank',
    'safe_anchor'
]

# Drop from X_train and X_test
df_reduced = df_reduced.drop(columns=features_to_drop)

# Update the final_features_vif list
final_features_vif = [feature for feature in final_features_vif if feature not
↪in features_to_drop]

# Add the newly engineered features

```

```

new_engineered_features = ['url_complexity', 'tag_to_link_ratio',
↪ 'domain_numeric_intensity', 'path_word_complexity']
final_features_vif.extend(new_engineered_features)

# Check final feature count
print("Total final features after update:", len(final_features_vif))
print("Updated Features:\n", final_features_vif)

```

Total final features after update: 27

Updated Features:

```

['prefix_suffix', 'nb_hyphens', 'domain_registration_length',
'ratio_digits_host', 'avg_word_path', 'ip', 'length_hostname', 'nb_www',
'empty_title', 'ratio_extRedirection', 'tld_in_subdomain', 'nb_hyperlinks',
'nb_and', 'domain_in_title', 'shortest_word_path', 'longest_words_raw',
'ratio_digits_url', 'shortest_word_host', 'links_in_tags', 'domain_age',
'phish_hints', 'ratio_extHyperlinks', 'nb_qm', 'url_complexity',
'tag_to_link_ratio', 'domain_numeric_intensity', 'path_word_complexity']

```

5.1 Insights and Recommendations

5.1.1 Key Insights from Feature Selection Process

The feature selection pipeline combined statistical rigor and machine learning techniques to ensure an optimal set of predictive variables:

1. Correlation Analysis

- Identified and removed highly correlated features ($\text{corr} > 0.9$) to reduce redundancy.
- **Dropped:** 'nb_eq', 'longest_word_path'
- Reduced feature count from **88 to 86**.

2. ANOVA F-Test (f_classif)

- Used to select the **top 30 features** based on **univariate analysis**.
- Suitable for identifying strong relationships between **numerical features** and the **categorical target**.

3. Random Forest Feature Importance

- Leveraged `feature_importances_` from a **trained Random Forest** to extract **top 30 influential features**.

4. Recursive Feature Elimination (RFE)

- Applied **RFE with Random Forest** as the estimator.
- Selected another **top 30 features**, enhancing robustness.

5. Feature Intersection (Union Strategy)

- Took the **intersection** of features selected by both **f__classif** and **RFE**.
- Resulted in a **robust and refined feature set** based on two complementary methods.

6. Variance Inflation Factor (VIF)

- Dropped **6 features** with **VIF > 10** to mitigate multicollinearity issues:
 - `length_words_raw`, `avg_words_raw`, `length_url`, `ratio_intHyperlinks`, `nb_slash`, `nb_dots`

5.1.2 Engineered Features That Add High Predictive Value

The following features were engineered to capture phishing-specific patterns:

Feature Name	Insight
<code>url_complexity</code>	Measures obfuscation via special characters in the URL. High values are often seen in phishing.
<code>tag_to_link_ratio</code>	Captures disproportionate script embedding relative to visible hyperlinks.
<code>domain_numeric_intensity</code>	Reflects digit-heavy domains with short registration times—typical of fraudulent domains.
<code>path_word_complexity</code>	Combines average and maximum path word lengths—phishing URLs often embed deep, confusing paths.

5.1.3 Dropped Redundant / Low-Predictive Features (Post-VIF)

The following features were removed to reduce redundancy or due to weak contribution:

- `domain_with_copyright`
- `ratio_intMedia`
- `google_index`
- `page_rank`
- `safe_anchor`
- (+6 VIF drops): `length_words_raw`, `avg_words_raw`, `length_url`, `ratio_intHyperlinks`, `nb_slash`, `nb_dots`

5.1.4 Final Recommendations

- Continue including engineered features in future model pipelines for domain-specific performance gains.

- Reapply VIF and correlation checks for each new dataset to ensure stability.
- Leverage tree-based models like **XGBoost** or **Random Forest** for feature importance validation.
- Normalize highly skewed features using **PowerTransformer** with `method='yeo-johnson'` to maintain model interpretability and performance.
- Consider permutation importance and SHAP values for model explainability.

5.2 Split Dataset into Train and Test set

```
[48]: from sklearn.model_selection import train_test_split

# Define final feature set and target
X_final = df_reduced[final_features_vif]
y_final = df_reduced['status']

# Perform stratified train-test split
X_train, X_test, y_train, y_test = train_test_split(
    X_final, y_final,
    test_size=0.2,
    random_state=42,
    stratify=y_final # maintain class distribution
)

# Generate report
train_size = X_train.shape[0]
test_size = X_test.shape[0]
total_size = len(y_final)

train_percent = round((train_size / total_size) * 100, 2)
test_percent = round((test_size / total_size) * 100, 2)

print(" Data Splitting Report:")
print(f" Total records: {total_size}")
print(f" Training set: {train_size} records ({train_percent}%)")
print(f" Testing set: {test_size} records ({test_percent}%)")

print("\n Target Distribution Check:")
print("Train set distribution:")
print(y_train.value_counts(normalize=True).map(lambda x: f"{x:.2%}"))

print("\nTest set distribution:")
print(y_test.value_counts(normalize=True).map(lambda x: f"{x:.2%}"))
```

```
Data Splitting Report:
Total records: 11430
Training set: 9144 records (80.0%)
Testing set: 2286 records (20.0%)
```

```
Target Distribution Check:
Train set distribution:
status
0    50.00%
1    50.00%
Name: proportion, dtype: object
```

```
Test set distribution:
status
1    50.00%
0    50.00%
Name: proportion, dtype: object
```

```
[49]: print("\nSkewness of Features:")
      X_train.skew()
```

Skewness of Features:

```
[49]: prefix_suffix      1.483091
      nb_hyphens          4.034987
      domain_registration_length 10.801880
      ratio_digits_host   5.615369
      avg_word_path       12.714639
      ip                  1.972296
      length_hostname     4.522406
      nb_www              0.264874
      empty_title         2.265138
      ratio_extRedirection 2.232868
      tld_in_subdomain    4.147150
      nb_hyperlinks       7.816814
      nb_and              10.090766
      domain_in_title     -1.328934
      shortest_word_path   4.649295
      longest_words_raw    14.463195
      ratio_digits_url     2.205006
      shortest_word_host   2.296740
      links_in_tags       -0.148617
      domain_age          0.168107
      phish_hints         3.249916
      ratio_extHyperlinks  1.018971
      nb_qm               2.480994
      url_complexity       4.126829
      tag_to_link_ratio    5.024884
      domain_numeric_intensity 5.877711
      path_word_complexity 32.492235
      dtype: float64
```

5.2.1 Handle Skewness

5.3 Skewness Handling Report

5.3.1 Technique Applied

- **Transformer:** Yeo–Johnson PowerTransformer
- **Library:** `sklearn.preprocessing.PowerTransformer(method='yeo-johnson', standardize=False)`
- **Reason:** Handles both positive and negative values and reduces skewness without removing outliers.

```
[50]: from sklearn.preprocessing import PowerTransformer

# Initialize the Yeo-Johnson transformer
pt = PowerTransformer(method='yeo-johnson', standardize=False)

# Fit the transformer on the training data and transform the training data
X_train_transformed = pt.fit_transform(X_train)

# Use the fitted transformer to transform the test data
X_test_transformed = pt.transform(X_test)

# Optional: Check skewness on transformed data
print("Skewness after Yeo-Johnson transform (Train):\n", pd.
      ↪DataFrame(X_train_transformed, columns=X_train.columns).skew().
      ↪sort_values(ascending=False))
print("Skewness after Yeo-Johnson transform (Test):\n", pd.
      ↪DataFrame(X_test_transformed, columns=X_test.columns).skew().
      ↪sort_values(ascending=False))
```

Skewness after Yeo-Johnson transform (Train):

tld_in_subdomain	4.147150
nb_and	3.512775
empty_title	2.265138
ratio_digits_host	2.199596
nb_qm	2.130345
ip	1.972296
phish_hints	1.701765
prefix_suffix	1.483091
ratio_digits_url	0.720100
domain_numeric_intensity	0.656725
ratio_extRedirection	0.650762
nb_hyphens	0.563168
tag_to_link_ratio	0.364356
ratio_extHyperlinks	0.319543

```

nb_www          0.219986
url_complexity  0.070897
shortest_word_host 0.018237
shortest_word_path 0.005782
avg_word_path   -0.013200
path_word_complexity -0.015818
length_hostname -0.031823
nb_hyperlinks   -0.040903
domain_registration_length -0.071173
longest_words_raw -0.097140
links_in_tags   -0.491997
domain_age      -0.765253
domain_in_title -1.328934
dtype: float64
Skewness after Yeo-Johnson transform (Test):
  tld_in_subdomain    4.035637
nb_and               3.346951
empty_title          2.298328
ratio_digits_host    2.277310
nb_qm                2.097703
ip                  1.886440
phish_hints          1.594152
prefix_suffix        1.474563
domain_numeric_intensity 0.676689
ratio_digits_url     0.643409
nb_hyphens           0.635152
ratio_extRedirection 0.630630
tag_to_link_ratio    0.355226
ratio_extHyperlinks  0.295872
nb_www              0.222480
length_hostname      0.205641
url_complexity       0.125121
avg_word_path        0.053828
path_word_complexity 0.034095
longest_words_raw    0.014290
shortest_word_path   0.004271
shortest_word_host   -0.036910
nb_hyperlinks        -0.082770
domain_registration_length -0.107386
links_in_tags        -0.507143
domain_age           -0.760209
domain_in_title      -1.301086
dtype: float64

```

- After Yeo-Johnson transformation, **most features' skewness** is reduced **close to zero**, indicating more symmetric distributions.
- This makes subsequent **scaling** (RobustScaler) and **model training** more stable and effective.

5.4 Scaling : RobustScaler()

```
[51]: from sklearn.preprocessing import RobustScaler

scaler = RobustScaler()
X_train_scaled = scaler.fit_transform(X_train_transformed)
X_test_scaled = scaler.transform(X_test_transformed)
```

6 Normalization/Scaling Report

6.1 Techniques Used:

- **Scaling Method Applied: RobustScaler**
 - **Reason for Selection:**
 - **RobustScaler** was chosen because it is robust to outliers. Unlike **StandardScaler** or **MinMaxScaler**, it scales features using **median** and **IQR (Interquartile Range)**, making it suitable for datasets with outliers, which is common in real-world data.
 - It helps ensure that features are on a similar scale, which is important for machine learning models like **SVM**, **Logistic Regression**, and **KNN**, which are sensitive to the scale of data.
-

6.2 Description of RobustScaler:

- **Scaler Formula:**

$$\text{scaled} = \frac{X - \text{median}(X)}{\text{IQR}(X)}$$

- ****Median:**** The middle value, less affected by outliers.
- ****IQR:**** The difference between the 75th and 25th percentiles, representing the range within

- **Impact of RobustScaler:**
 - **Prevents Outlier Influence:** The scaling technique is **not influenced by extreme values**.
 - **Preserves Distribution:** Data is centered and scaled based on the distribution within the interquartile range, making it **robust to skewed distributions**.
-

```
[52]: # Calculate original distribution (min, max)
original_stats = X_train.agg(['min', 'max']).T
original_stats.columns = ['Original Min', 'Original Max']

X_train_scaled_df = pd.DataFrame(X_train_scaled, columns=X_train.columns)

# Calculate scaled distribution (min, max)
```

```

scaled_stats = X_train_scaled_df.agg(['min', 'max']).T
scaled_stats.columns = ['Scaled Min', 'Scaled Max']

# Combine both into a single table for comparison
comparison_df = pd.concat([original_stats, scaled_stats], axis=1)

# Print results
print("Before-and-After Feature Scaling (RobustScaler):\n")
print(comparison_df.round(3))

```

Before-and-After Feature Scaling (RobustScaler):

	Original Min	Original Max	Scaled Min	Scaled Max
prefix_suffix	0.0	1.000	-0.000	0.145
nb_hyphens	0.0	32.000	-0.000	1.563
domain_registration_length	-1.0	29829.000	-2.223	5.435
ratio_digits_host	0.0	0.800	-0.000	0.024
avg_word_path	0.0	206.000	-0.840	2.999
ip	0.0	1.000	-0.000	0.104
length_hostname	4.0	214.000	-3.638	4.130
nb_www	0.0	2.000	-0.000	1.342
empty_title	0.0	1.000	-0.000	0.087
ratio_extRedirection	0.0	2.000	-0.000	1.465
tld_in_subdomain	0.0	1.000	-0.000	0.034
nb_hyperlinks	0.0	4659.000	-1.230	2.729
nb_and	0.0	19.000	-0.000	0.075
domain_in_title	0.0	1.000	-10.750	0.000
shortest_word_path	0.0	40.000	-0.797	1.758
longest_words_raw	2.0	829.000	-3.997	3.338
ratio_digits_url	0.0	0.724	-0.000	1.486
shortest_word_host	1.0	39.000	-1.885	2.521
links_in_tags	0.0	100.000	-0.784	0.216
domain_age	-12.0	12874.000	-2.338	0.862
phish_hints	0.0	10.000	-0.000	0.182
ratio_extHyperlinks	0.0	1.000	-0.432	0.885
nb_qm	0.0	3.000	-0.000	0.096
url_complexity	0.0	34.000	-0.681	1.447
tag_to_link_ratio	0.0	50.000	-0.507	1.067
domain_numeric_intensity	-0.8	3828.649	-1.659	0.927
path_word_complexity	0.0	83636.000	-0.851	2.596

6.3 Before-and-After Comparison of Numerical Feature Distributions:

6.3.1 Before Scaling:

- Features can have **different ranges** (e.g., one feature ranges from 0 to 10, while another ranges from 100 to 1000).
- Outliers could heavily influence the distributions (e.g., extremely large values may shift the

mean).

6.3.2 After Scaling (RobustScaler):

- Features are scaled within a similar range but **without the influence of outliers**.
- The **central tendency** (median) and **spread** (IQR) are preserved and adjusted for each feature, so all features are on a comparable scale for model training.

All feature values are now on a similar scale centered around 0, making the model training more stable and faster.

```
[53]: # Final split dataset ready for model training
X_train_scaled
X_test_scaled
```

```
[53]: array([[ 0.1450004 ,  1.          , -1.97451712, ..., -0.50692832,
        -0.          ,  0.14603119],
       [ 0.1450004 ,  1.          , -0.17208061, ...,  0.349175   ,
        -0.          ,  0.59878878],
       [ 0.1450004 ,  1.          , -0.03154315, ..., -0.50692832,
        -0.          , -0.02040373],
       ...,
       [-0.          , -0.          ,  2.09673827, ..., -0.20355885,
        -0.          , -0.85079951],
       [-0.          ,  1.50903381,  0.09939827, ...,  0.35561974,
        -0.          , -0.00963592],
       [ 0.1450004 ,  1.          , -1.97451712, ..., -0.42727707,
        -0.          , -0.85079951]], shape=(2286, 27))
```

```
[ ]:
```